



Ngaire44 / PROJECT

[Code](#)[Issues](#)[Pull requests](#)[Actions](#)[Projects](#)[Wiki](#)[Security](#)[PROJECT / notebook.ipynb](#)

Ngaire44 final project

180d5b7 · 22 minutes ago



5575 lines (5575 loc) · 887 KB

# SyriaTel Customer Churn Project

## Business Understanding

SyriaTel, a telecommunications company, experiences customer churn at around 15%, meaning thousands of customers leave every year. This results in lost revenue, increased acquisition costs, and weakened customer loyalty. While churn prediction is a common use case, we aim to go beyond prediction by identifying the key behavioral and service-related factors driving churn. This will help SyriaTel target interventions effectively and reduce customer attrition.

This project aims:

- To build a predictive model that identifies customers at high risk of churn
- To identify key factors likely to make a customer churn; service quality, call charges, customer support
- To recommend actions that will reduce churn
- 5% reduction in churn rate through targeted retention strategies

Incase this dataset is imbalanced, accuracy alone is not a good measure. Instead, we will evaluate success using:

- Recall  $\geq 85$  (Sensitivity) for churn class - ensure the model identifies as many churners as possible, since missing them means lost revenue.
- Accuracy  $\geq 85$  so predictions are overall correct
- F1-score - balance between Precision and Recall.
- AUC  $\geq 0.85$  -overall ability of the model to discriminate churners vs. non-churners.

## Data Understanding

### Columns description

- state: The U.S. state where the customer resides
- account length: Number of days the customer has had an account with SyriaTel
- area code: The customer's telephone area code
- phone number: The unique phone number assigned to the customer.
- international plan: A binary indicator (yes/no) of whether the customer has subscribed to an international calling plan.
- voice mail plan: A binary indicator (yes/no) of whether the customer has subscribed to a voicemail plan.
- number vmail messages: The number of voicemail messages recorded by the customer.
- total day minutes: The total number of minutes the customer used during

daytime calls

- total day calls: The total number of calls made during the day.
- total day charge: The total charges incurred from daytime calls, directly proportional to total day minutes.
- total eve minutes: The total number of minutes the customer used during evening calls
- total eve calls: The total number of calls made during the evening.
- total eve charge: The total charges incurred from evening calls, directly proportional to total eve minutes.
- total night minutes: The total number of minutes the customer used during nighttime calls
- total night calls: The total number of calls made during the night.
- total night charge: The total charges incurred from nighttime calls, directly proportional to total night minutes.
- total intl minutes: The total number of minutes spent on international calls.
- total intl calls: The total number of international calls made by the customer.
- total intl charge: The total charges incurred from international calls, directly proportional to total intl minutes.
- customer service calls: The number of calls the customer made to customer service.
- churn: indicates whether the customer has churned (yes/no); Target variable

In [1]:

```
# import libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')

from scipy.stats import chi2_contingency, ttest_ind
from xgboost import XGBClassifier

from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import f1_score, roc_auc_score, roc_curve, \
    classification_report, confusion_matrix, accuracy_score, \
    recall_score, precision_score, auc, make_scorer, average_precision_score
from sklearn.model_selection import train_test_split, GridSearchCV, \
    cross_validate, StratifiedKFold, StratifiedShuffleSplit
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.calibration import CalibratedClassifierCV
from sklearn.inspection import permutation_importance
from imblearn.over_sampling import SMOTE
```

In [2]:

```
#load datasets
df = pd.read_csv('bigml.csv')
df.head()
```

Out[2]:

state	account length	area code	phone number	international plan	voice mail plan	number vmail plan	total day messages	total day minutes	total day calls
-------	----------------	-----------	--------------	--------------------	-----------------	-------------------	--------------------	-------------------	-----------------

0	KS	128	415	382-4657	no	yes	25	265.1	110
1	OH	107	415	371-7191	no	yes	26	161.6	123
2	NJ	137	415	358-1921	no	no	0	243.4	114
3	OH	84	408	375-9999	yes	no	0	299.4	71
4	OK	75	415	330-6626	yes	no	0	166.7	113

5 rows × 21 columns

In [3]:

```
df.tail()
```

Out[3]:

	state	account length	area code	phone number	international plan	voice mail plan	number vmail messages	total day minutes	t c
3328	AZ	192	415	414-4276	no	yes	36	156.2	
3329	WV	68	415	370-3271	no	no	0	231.1	
3330	RI	28	510	328-8230	no	no	0	180.8	
3331	CT	184	510	364-6381	yes	no	0	213.8	
3332	TN	74	415	400-4344	no	yes	25	234.4	

5 rows × 21 columns

observations: the columns are uniform

In [4]:

```
#shape of the dataset
print(f"The dataset has {df.shape[0]} rows and {df.shape[1]} columns")
```

The dataset has 3333 rows and 21 columns

In [5]:

```
# investigate the data types
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3333 entries, 0 to 3332
Data columns (total 21 columns):
 #   Column           Non-Null Count  Dtype  
 ---  -- 
 0   state            3333 non-null   object 
 1   account length  3333 non-null   int64  
 2   area code        3333 non-null   int64  
 3   phone number    3333 non-null   object 
 4   international plan 3333 non-null   object 
 5   voice mail plan 3333 non-null   object 
```

```

6    number vmail messages    3333 non-null    int64
7    total day minutes     3333 non-null    float64
8    total day calls       3333 non-null    int64
9    total day charge      3333 non-null    float64
10   total eve minutes     3333 non-null    float64
11   total eve calls       3333 non-null    int64
12   total eve charge      3333 non-null    float64
13   total night minutes   3333 non-null    float64
14   total night calls     3333 non-null    int64
15   total night charge    3333 non-null    float64
16   total intl minutes    3333 non-null    float64
17   total intl calls      3333 non-null    int64
18   total intl charge     3333 non-null    float64
19   customer service calls 3333 non-null    int64
20   churn                  3333 non-null    bool
dtypes: bool(1), float64(8), int64(8), object(4)
memory usage: 524.2+ KB

```

observations:

- We will one hot encode the intl plan, voicemail plan and churn
- The dataset has no missing values

In [6]:

```
#stat summary
df.describe()
```

Out[6]:

	account length	area code	number vmail messages	total day minutes	total day calls	
count	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000	333
mean	101.064806	437.182418	8.099010	179.775098	100.435644	3
std	39.822106	42.371290	13.688365	54.467389	20.069084	
min	1.000000	408.000000	0.000000	0.000000	0.000000	
25%	74.000000	408.000000	0.000000	143.700000	87.000000	2
50%	101.000000	415.000000	0.000000	179.400000	101.000000	3
75%	127.000000	510.000000	20.000000	216.400000	114.000000	3
max	243.000000	510.000000	51.000000	350.800000	165.000000	5

In [7]:

```
#unique values
for col in df:
    uniq_values = df[col].unique()
    print(f'{col}\n {df[col].unique()}\n')
```

```

state
['KS' 'OH' 'NJ' 'OK' 'AL' 'MA' 'MO' 'LA' 'WV' 'IN' 'RI' 'IA' 'MT' 'NY'
 'ID' 'VT' 'VA' 'TX' 'FL' 'CO' 'AZ' 'SC' 'NE' 'WY' 'HI' 'IL' 'NH' 'GA'
 'AK' 'MD' 'AR' 'WI' 'OR' 'MI' 'DE' 'UT' 'CA' 'MN' 'SD' 'NC' 'WA' 'NM'
 'NV' 'DC' 'KY' 'ME' 'MS' 'TN' 'PA' 'CT' 'ND']

account length
[128 107 137 84 75 118 121 147 117 141 65 74 168 95 62 161 85 9
3
 76 73 77 130 111 132 174 57 54 20 49 142 172 12 72 36 78 136
149 98 135 34 160 64 59 119 97 52 60 10 96 87 81 68 125 116

```

```

38 40 43 113 126 150 138 162 90 50 82 144 46 70 55 106 94 155
80 104 99 120 108 122 157 103 63 112 41 193 61 92 131 163 91 127
110 140 83 145 56 151 139 6 115 146 185 148 32 25 179 67 19 170
164 51 208 53 105 66 86 35 88 123 45 100 215 22 33 114 24 101
143 48 71 167 89 199 166 158 196 209 16 39 173 129 44 79 31 124
37 159 194 154 21 133 224 58 11 109 102 165 18 30 176 47 190 152
26 69 186 171 28 153 169 13 27 3 42 189 156 134 243 23 1 205
200 5 9 178 181 182 217 177 210 29 180 2 17 7 212 232 192 195
197 225 184 191 201 15 183 202 8 175 4 188 204 221]

```

area code

```
[415 408 510]
```

phone number

```
['382-4657' '371-7191' '358-1921' ... '328-8230' '364-6381' '400-4344']
```

international plan

```
['no' 'yes']
```

voice mail plan

```
['yes' 'no']
```

number vmail messages

```
[25 26 0 24 37 27 33 39 30 41 28 34 46 29 35 21 32 42 36 22 23 43 31 3
8
40 48 18 17 45 16 20 14 19 51 15 11 12 47 8 44 49 4 10 13 50 9]
```

total day minutes

```
[265.1 161.6 243.4 ... 321.1 231.1 180.8]
```

total day calls

```
[110 123 114 71 113 98 88 79 97 84 137 127 96 70 67 139 66 9
0
117 89 112 103 86 76 115 73 109 95 105 121 118 94 80 128 64 106
102 85 82 77 120 133 135 108 57 83 129 91 92 74 93 101 146 72
99 104 125 61 100 87 131 65 124 119 52 68 107 47 116 151 126 122
111 145 78 136 140 148 81 55 69 158 134 130 63 53 75 141 163 59
132 138 54 58 62 144 143 147 36 40 150 56 51 165 30 48 60 42
0 45 160 149 152 142 156 35 49 157 44]
```

total day charge

```
[45.07 27.47 41.38 ... 54.59 39.29 30.74]
```

total eve minutes

```
[197.4 195.5 121.2 ... 153.4 288.8 265.9]
```

total eve calls

```
[ 99 103 110 88 122 101 108 94 80 111 83 148 71 75 76 97 90 6
5
93 121 102 72 112 100 84 109 63 107 115 119 116 92 85 98 118 74
117 58 96 66 67 62 77 164 126 142 64 104 79 95 86 105 81 113
106 59 48 82 87 123 114 140 128 60 78 125 91 46 138 129 89 133
136 57 135 139 51 70 151 137 134 73 152 168 68 120 69 127 132 143
61 124 42 54 131 52 149 56 37 130 49 146 147 55 12 50 157 155
45 144 36 156 53 141 44 153 154 150 43 0 145 159 170]
```

total eve charge

```
[16.78 16.62 10.3 ... 13.04 24.55 22.6 ]
```

total night minutes

```
[244.7 254.4 162.6 ... 280.9 120.1 279.1]
```

total night calls

```
[ 91 103 104 89 121 118 96 90 97 111 94 128 115 99 75 108 74 13
5
```

64	78	105	68	102	148	98	116	71	109	107	135	92	86	127	79	87	129
57	77	95	54	106	53	67	139	60	100	61	73	113	76	119	88	84	62
137	72	142	114	126	122	81	123	117	82	80	120	130	134	59	112	132	110
101	150	69	131	83	93	124	136	125	66	143	58	55	85	56	70	46	42
152	44	145	50	153	49	175	63	138	154	140	141	146	65	51	151	158	155
157	147	144	149	166	52	33	156	38	36	48	164]						

total night charge

[11.01	11.45	7.32	8.86	8.41	9.18	9.57	9.53	9.71	14.69	9.4	8.8						
2																	
6.35	8.65	9.14	7.23	4.02	5.83	7.46	8.68	9.43	8.18	8.53	10.67						
11.28	8.22	4.59	8.17	8.04	11.27	11.08	13.2	12.61	9.61	6.88	5.82						
10.25	4.58	8.47	8.45	5.5	14.02	8.03	11.94	7.34	6.06	10.9	6.44						
3.18	10.66	11.21	12.73	10.28	12.16	6.34	8.15	5.84	8.52	7.5	7.48						
6.21	11.95	7.15	9.63	7.1	6.91	6.69	13.29	11.46	7.76	6.86	8.16						
12.15	7.79	7.99	10.29	10.08	12.53	7.91	10.02	8.61	14.54	8.21	9.09						
4.93	11.39	11.88	5.75	7.83	8.59	7.52	12.38	7.21	5.81	8.1	11.04						
11.19	8.55	8.42	9.76	9.87	10.86	5.36	10.03	11.15	9.51	6.22	2.59						
7.65	6.45	9.	6.4	9.94	5.08	10.23	11.36	6.97	10.16	7.88	11.91						
6.61	11.55	11.76	9.27	9.29	11.12	10.69	8.8	11.85	7.14	8.71	11.42						
4.94	9.02	11.22	4.97	9.15	5.45	7.27	12.91	7.75	13.46	6.32	12.13						
11.97	6.93	11.66	7.42	6.19	11.41	10.33	10.65	11.92	4.77	4.38	7.41						
12.1	7.69	8.78	9.36	9.05	12.7	6.16	6.05	10.85	8.93	3.48	10.4						
5.05	10.71	9.37	6.75	8.12	11.77	11.49	11.06	11.25	11.03	10.82	8.91						
8.57	8.09	10.05	11.7	10.17	8.74	5.51	11.11	3.29	10.13	6.8	8.49						
9.55	11.02	9.91	7.84	10.62	9.97	3.44	7.35	9.79	8.89	8.14	6.94						
10.49	10.57	10.2	6.29	8.79	10.04	12.41	15.97	9.1	11.78	12.75	11.07						
12.56	8.63	8.02	10.42	8.7	9.98	7.62	8.33	6.59	13.12	10.46	6.63						
8.32	9.04	9.28	10.76	9.64	11.44	6.48	10.81	12.66	11.34	8.75	13.05						
11.48	14.04	13.47	5.63	6.6	9.72	11.68	6.41	9.32	12.95	13.37	9.62						
6.03	8.25	8.26	11.96	9.9	9.23	5.58	7.22	6.64	12.29	12.93	11.32						
6.85	8.88	7.03	8.48	3.59	5.86	6.23	7.61	7.66	13.63	7.9	11.82						
7.47	6.08	8.4	5.74	10.94	10.35	10.68	4.34	8.73	5.14	8.24	9.99						
13.93	8.64	11.43	5.79	9.2	10.14	12.11	7.53	12.46	8.46	8.95	9.84						
10.8	11.23	10.15	9.21	14.46	6.67	12.83	9.66	9.59	10.48	8.36	4.84						
10.54	8.39	7.43	9.06	8.94	11.13	8.87	8.5	7.6	10.73	9.56	10.77						
7.73	3.47	11.86	8.11	9.78	9.42	9.65	7.	7.39	9.88	6.56	5.92						
6.95	15.71	8.06	4.86	7.8	8.58	10.06	5.21	6.92	6.15	13.49	9.38						
12.62	12.26	8.19	11.65	11.62	10.83	7.92	7.33	13.01	13.26	12.22	11.58						
5.97	10.99	8.38	9.17	8.08	5.71	3.41	12.63	11.79	12.96	7.64	6.58						
10.84	10.22	6.52	5.55	7.63	5.11	5.89	10.78	3.05	11.89	8.97	10.44						
10.5	9.35	5.66	11.09	9.83	5.44	10.11	6.39	11.93	8.62	12.06	6.02						
8.85	5.25	8.66	6.73	10.21	11.59	13.87	7.77	10.39	5.54	6.62	13.33						
6.24	12.59	6.3	6.79	8.28	9.03	8.07	5.52	12.14	10.59	7.54	7.67						
5.47	8.81	8.51	13.45	8.77	6.43	12.01	12.08	7.07	6.51	6.84	9.48						
13.78	11.54	11.67	8.13	10.79	7.13	4.72	4.64	8.96	13.03	6.07	3.51						
6.83	6.12	9.31	9.58	4.68	5.32	9.26	11.52	9.11	10.55	11.47	9.3						
13.82	8.44	5.77	10.96	11.74	8.9	10.47	7.85	10.92	4.74	9.74	10.43						
9.96	10.18	9.54	7.89	12.36	8.54	10.07	9.46	7.3	11.16	9.16	10.19						
5.99	10.88	5.8	7.19	4.55	8.31	8.01	14.43	8.3	14.3	6.53	8.2						
11.31	13.	6.42	4.24	7.44	7.51	13.1	9.49	6.14	8.76	6.65	10.56						
6.72	8.29	12.09	5.39	2.96	7.59	7.24	4.28	9.7	8.83	13.3	11.37						
9.33	5.01	3.26	11.71	8.43	9.68	15.56	9.8	3.61	6.96	11.61	12.81						
10.87	13.84	5.03	5.17	2.03	10.34	9.34	7.95	10.09	9.95	7.11	9.22						
6.13	11.05	9.89	9.39	14.06	10.26	13.31	15.43	16.39	6.27	10.64	11.5						
12.48	8.27	13.53	10.36	12.24	8.69	10.52	9.07	11.51	9.25	8.72	6.78						
8.6	11.84	5.78	5.85	12.3	5.76	12.07	9.6	8.84	12.39	10.1	9.73						
2.85	6.66	2.45	5.28	11.73	10.75	7.74	6.76	6.	7.58	13.69	7.93						
7.68	9.75	4.96	5.49	11.83	7.18	9.19	7.7	7.25	10.74	4.27	13.8						
9.12	4.75	7.78	11.63	7.55	2.25	9.45	9.86	7.71	4.95	7.4	11.17						
11.33	6.82	13.7	1.97	10.89	12.77	10.31	5.23	5.27	9.41	6.09	10.61						
7.29	4.23	7.57	3.67	12.69	14.5	5.95	7.87	5.96	5.94	12.23	4.9						
12.33	6.89	9.67	12.68	12.87	3.7	6.04	13.13	15.74	11.87	4.7	4.67						
7.05	5.42	4.09	5.73	9.47	8.05	6.87	3.71	15.86	7.49	11.69	6.46						

10.45	12.9	5.41	11.26	1.04	6.49	6.37	12.21	6.77	12.65	7.86	9.44
4.3	7.38	5.02	10.63	2.86	17.19	8.67	8.37	6.9	10.93	10.38	7.36
10.27	10.95	6.11	4.45	11.9	15.01	12.84	7.45	6.98	11.72	7.56	11.38
10.	10.	6.01	5.00	6.01	10.00	10.00	10.00	6.00	11.00	6.00	7.00

## PROJECT / notebook.ipynb

[Top](#)[Preview](#)[Code](#)[Blame](#)[Raw](#)[Edit](#)

```

5.57 3.94 4.41 13.27 10.24 4.25 12.89 5.72 12.5 11.29 3.25 11.53
9.82 7.26 4.1 10.37 4.98 6.74 12.52 14.56 8.34 3.82 3.86 13.97
11.57 6.5 13.58 14.32 13.75 11.14 14.18 9.13 4.46 4.83 9.69 14.13
7.16 7.98 13.66 14.78 11.2 9.93 11. 5.29 9.92 4.29 11.1 10.51
12.49 4.04 12.94 7.09 6.71 7.94 5.31 5.98 7.2 14.82 13.21 12.32
10.58 4.92 6.2 4.47 11.98 6.18 7.81 4.54 5.37 7.17 5.33 14.1
5.7 12.18 8.98 5.1 14.67 13.95 16.55 11.18 4.44 4.73 2.55 6.31
2.43 9.24 7.37 13.42 12.42 11.8 14.45 2.89 13.23 12.6 13.18 12.19
14.81 6.55 11.3 12.27 13.98 8.23 15.49 6.47 13.48 13.59 13.25 17.77
13.9 3.97 11.56 14.08 13.6 6.26 4.61 12.76 15.76 6.38 3.6 12.8
5.9 7.97 5. 10.97 5.88 12.34 12.03 14.97 15.06 12.85 6.54 11.24
12.64 7.06 5.38 13.14 3.99 3.32 4.51 4.12 3.93 2.4 11.75 4.03
15.85 6.81 14.25 14.09 16.42 6.7 12.74 2.76 12.12 6.99 6.68 11.81
7.96 5.06 13.16 2.13 13.17 5.12 5.65 12.37 10.53]

```

**total intl minutes**

```

[10. 13.7 12.2 6.6 10.1 6.3 7.5 7.1 8.7 11.2 12.7 9.1 12.3 13.1
 5.4 13.8 8.1 13. 10.6 5.7 9.5 7.7 10.3 15.5 14.7 11.1 14.2 12.6
 11.8 8.3 14.5 10.5 9.4 14.6 9.2 3.5 8.5 13.2 7.4 8.8 11. 7.8
 6.8 11.4 9.3 9.7 10.2 8. 5.8 12.1 12. 11.6 8.2 6.2 7.3 6.1
 11.7 15. 9.8 12.4 8.6 10.9 13.9 8.9 7.9 5.3 4.4 12.5 11.3 9.
 9.6 13.3 20. 7.2 6.4 14.1 14.3 6.9 11.5 15.8 12.8 16.2 0. 11.9
 9.9 8.4 10.8 13.4 10.7 17.6 4.7 2.7 13.5 12.9 14.4 10.4 6.7 15.4
 4.5 6.5 15.6 5.9 18.9 7.6 5. 7. 14. 18. 16. 14.8 3.7 2.
 4.8 15.3 6. 13.6 17.2 17.5 5.6 18.2 3.6 16.5 4.6 5.1 4.1 16.3
 14.9 16.4 16.7 1.3 15.2 15.1 15.9 5.5 16.1 4. 16.9 5.2 4.2 15.7
 17. 3.9 3.8 2.2 17.1 4.9 17.9 17.3 18.4 17.8 4.3 2.9 3.1 3.3
 2.6 3.4 1.1 18.3 16.6 2.1 2.4 2.5]

```

**total intl calls**

```
[ 3 5 7 6 4 2 9 19 1 10 15 8 11 0 12 13 18 14 16 20 17]
```

**total intl charge**

```

[2.7 3.7 3.29 1.78 2.73 1.7 2.03 1.92 2.35 3.02 3.43 2.46 3.32 3.54
 1.46 3.73 2.19 3.51 2.86 1.54 2.57 2.08 2.78 4.19 3.97 3. 3.83 3.4
 3.19 2.24 3.92 2.84 2.54 3.94 2.48 0.95 2.3 3.56 2. 2.38 2.97 2.11
 1.84 3.08 2.51 2.62 2.75 2.16 1.57 3.27 3.24 3.13 2.21 1.67 1.97 1.65
 3.16 4.05 2.65 3.35 2.32 2.94 3.75 2.4 2.13 1.43 1.19 3.38 3.05 2.43
 2.59 3.59 5.4 1.94 1.73 3.81 3.86 1.86 3.11 4.27 3.46 4.37 0. 3.21
 2.67 2.27 2.92 3.62 2.89 4.75 1.27 0.73 3.65 3.48 3.89 2.81 1.81 4.16
 1.22 1.76 4.21 1.59 5.1 2.05 1.35 1.89 3.78 4.86 4.32 4. 1. 0.54
 1.3 4.13 1.62 3.67 4.64 4.73 1.51 4.91 0.97 4.46 1.24 1.38 1.11 4.4
 4.02 4.43 4.51 0.35 4.1 4.08 4.29 1.49 4.35 1.08 4.56 1.4 1.13 4.24
 4.59 1.05 1.03 0.59 4.62 1.32 4.83 4.67 4.97 4.81 1.16 0.78 0.84 0.89
 0.7 0.92 0.3 4.94 4.48 0.57 0.65 0.68]

```

**customer service calls**

```
[1 0 2 3 4 5 7 9 6 8]
```

**churn**

```
[False True]
```

## Data Preparation

In [8]:

```
#make a copy
data = df.copy(deep=True)
```

In [9]:

```
#remove spaces in column names
data.columns = data.columns.str.replace(" ","_")
data.columns
```

Out[9]: Index(['state', 'account\_length', 'area\_code', 'phone\_number',  
 'international\_plan', 'voice\_mail\_plan', 'number\_vmail\_messages',  
 'total\_day\_minutes', 'total\_day\_calls', 'total\_day\_charge',  
 'total\_eve\_minutes', 'total\_eve\_calls', 'total\_eve\_charge',  
 'total\_night\_minutes', 'total\_night\_calls', 'total\_night\_charge',  
 'total\_intl\_minutes', 'total\_intl\_calls', 'total\_intl\_charge',  
 'customer\_service\_calls', 'churn'],  
 dtype='object')

In [10]:

```
#change churn data type to object
data.churn = data['churn'].astype('0')
```

In [11]:

```
#drop phone numbers
del data['phone_number']
data
```

Out[11]:

	state	account_length	area_code	international_plan	voice_mail_plan	nu
0	KS	128	415	no	yes	
1	OH	107	415	no	yes	
2	NJ	137	415	no	no	
3	OH	84	408	yes	no	
4	OK	75	415	yes	no	
...	...	...	...	...	...	...
3328	AZ	192	415	no	yes	
3329	WV	68	415	no	no	
3330	RI	28	510	no	no	
3331	CT	184	510	yes	no	
3332	TN	74	415	no	yes	

3333 rows × 20 columns

In [12]:

```
#check for duplicates,missing values
print(f"The dataset has {data.duplicated().sum()} duplicated values")
print(f"The dataset has: \n{data.isna().sum()} missing values")
```

The dataset has 0 duplicated values

The dataset has:

state	0
account_length	0

```

area_code          0
international_plan 0
voice_mail_plan    0
number_vmail_messages 0
total_day_minutes 0
total_day_calls    0
total_day_charge   0
total_eve_minutes  0
total_eve_calls    0
total_eve_charge   0
total_night_minutes 0
total_night_calls  0
total_night_charge 0
total_intl_minutes 0
total_intl_calls   0
total_intl_charge  0
customer_service_calls 0
churn             0
dtype: int64 missing values


```

In [13]:

```

#check for outliers
plt.figure(figsize=(15,8))
sns.boxplot(data,color='blue')
plt.xticks(rotation=45)
plt.tight_layout()

```



Observations: The outliers are genuine

- Call times may vary depending on clients' job, location or lifestyle
- Call charges are dependent on time spent on call
- Customer service calls may increase as a result of complaints from clients

In [14]:

```

#confirm data types
data.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3333 entries, 0 to 3332
Data columns (total 20 columns):
 #   Column           Non-Null Count  Dtype  
 ---  -- 
 0   state            3333 non-null   object 
 1   account_length   3333 non-null   int64 

```

```

2   area_code           3333 non-null    int64
3   international_plan 3333 non-null    object
4   voice_mail_plan    3333 non-null    object
5   number_vmail_messages 3333 non-null    int64
6   total_day_minutes  3333 non-null    float64
7   total_day_calls    3333 non-null    int64
8   total_day_charge   3333 non-null    float64
9   total_eve_minutes  3333 non-null    float64
10  total_eve_calls   3333 non-null    int64
11  total_eve_charge  3333 non-null    float64
12  total_night_minutes 3333 non-null    float64
13  total_night_calls  3333 non-null    int64
14  total_night_charge 3333 non-null    float64
15  total_intl_minutes 3333 non-null    float64
16  total_intl_calls   3333 non-null    int64
17  total_intl_charge  3333 non-null    float64
18  customer_service_calls 3333 non-null    int64
19  churn              3333 non-null    object
dtypes: float64(8), int64(8), object(4)
memory usage: 520.9+ KB


```

```
In [15]: #save the cleaned dataset
data.to_csv('clean_customer_churn.csv', index=False)
```

## Exploratory Data Analysis


### Univariate Analysis

```
In [16]: #check for imbalance in churn and plans
label = ['no', 'yes']
fig, (ax1, ax2, ax3) = plt.subplots(1, 3, figsize=(10, 5))
ax1.pie(data['churn'].value_counts(), autopct='%.2f%%', labels=label)
ax1.set_title("churn")
ax2.pie(data['international_plan'].value_counts(), autopct='%.2f%%', labels=label)
ax2.set_title("international plan")
ax3.pie(data['voice_mail_plan'].value_counts(), autopct='%.2f%%', labels=label)
ax3.set_title("voice mail plan")
plt.tight_layout();
```



```
In [17]: sns.histplot(data['customer_service_calls'], bins=10)
plt.title('customer service calls')
plt.tight_layout();
```






In [18]:

```
usage_cols = ['total_day_minutes', 'total_eve_minutes', 'total_night_minutes',
              'total_day_calls', 'total_eve_calls', 'total_night_calls',
              axes = axes.flatten()

for i,col in enumerate(usage_cols):
    sns.histplot(data[col], bins='auto', kde=True, ax=axes[i])
    axes[i].set_title(f"{col}Distribution")
    axes[i].set_xlabel(col)
    axes[i].set_ylabel('Frequency')
    plt.tight_layout();
```



In [19]:

```
data.columns
```

```
Out[19]: Index(['state', 'account_length', 'area_code', 'international_plan',
       'voice_mail_plan', 'number_vmail_messages', 'total_day_minutes',
       'total_day_calls', 'total_day_charge', 'total_eve_minutes',
       'total_eve_calls', 'total_eve_charge', 'total_night_minutes',
       'total_night_calls', 'total_night_charge', 'total_intl_minutes',
       'total_intl_calls', 'total_intl_charge', 'customer_service_calls'])
```

```
'churn'],
dtype='object')
```

## Observations

- 15% of the customers churn; we have an imbalance
- Most customers call twice while few call upto 8 times; Do those who call more than 6 times churn?
- The usage columns are normally distributed apart from total internation calls
- customer service calls column is skewed to the right
- intl plan and voicemail plan columns are imbalanced; we may use SMOTE

## Bivariate analysis

- We will compare features to churn

In [20]:

```
#H0: there is no association between customer plan and churn
# H1: there is an association between customer plan and churn

customer_plan = data[['international_plan', 'voice_mail_plan']]
chi = {}
for col in customer_plan:
    contingency = pd.crosstab(data[col], data['churn'])
    chi2, p_value, _, _ = chi2_contingency(contingency)
    chi[col] = p_value
print(f"p_value for international plan is {chi['international_plan']:.4f}\n\n")


#let's visualize to see the association
fig, axes = plt.subplots(1, 2, figsize=(12, 6))


# Loop through both subplots and columns
for ax, col in zip(axes, customer_plan.columns):
    # Get contingency table
    customer_plan_ct = pd.crosstab(data[col], data['churn'])
    customer_plan_ct.plot(kind='bar', ax=ax, legend=(col == 'voice_mail_plan'))
    churn_rate = (customer_plan_ct[True] / customer_plan_ct.sum(axis=1))
    # Add churn % labels on top of orange bars

    for i in range(len(customer_plan_ct)):
        height = customer_plan_ct.iloc[i][True]
        percent = churn_rate.iloc[i]
        ax.text(i, height + 10, f"{percent:.1f}%", ha='center', color='red')

    # Titles and axis labels
    ax.set_title(f"Churn by {col.replace('_', ' ')}.title()")
    ax.set_xlabel(col.replace('_', ' ').title())
    ax.set_ylabel("Number of Customers")
    ax.set_xticklabels(customer_plan_ct.index, rotation=0);
```

p\_value for international plan is 0.0000  
 p-value for voice mail plan is 0.0000





In [21]:

```
# H0: There is no significant difference in means of churn vs non-churn
# H1: There is significant difference in means of churn vs non-churn in these features

num_cols = ['account_length', 'total_day_minutes', 'total_day_calls', 'total_charge',
            'total_night_calls', 'total_intl_minutes', 'total_intl_calls',
            'total_eve_charge', 'total_night_charge', 'total_intl_charge']
anova = []
for col in num_cols:
    group0 = data[data['churn'] == 0][col]
    group1 = data[data['churn'] == 1][col]
    t_stat, p_value = ttest_ind(group0, group1, equal_var=False)
    is_significant = p_value < 0.05
    anova.append({'Feature': col, 'p-value': p_value, 'Significant (<0.05)': is_significant})

anova_df = pd.DataFrame(anova)
anova_df = anova_df.sort_values(by='p-value')
print(anova_df[anova_df['Significant (<0.05)']]])
```

	Feature	p-value	Significant (<0.05)
1	total_day_minutes	1.218968e-20	True
11	total_day_charge	1.219876e-20	True
9	customer_service_calls	5.270040e-18	True
10	number_vmail_messages	8.764782e-09	True
3	total_eve_minutes	1.839080e-07	True
12	total_eve_charge	1.842608e-07	True
14	total_intl_charge	9.025887e-05	True
7	total_intl_minutes	9.065715e-05	True
8	total_intl_calls	3.185777e-03	True
13	total_night_charge	3.027154e-02	True
5	total_night_minutes	3.028049e-02	True


In [22]:


```
# customer service by churn
service_ct = pd.crosstab(data['customer_service_calls'], data['churn'])
ax5 = service_ct.plot(kind='bar')

service_churn_rate = (service_ct[True] / service_ct.sum(axis=1)) * 100
service_bar_position = range(len(service_ct))
for i in service_bar_position:
    height_service = service_ct.iloc[i][True]
    service_percent = service_churn_rate.iloc[i]
    ax5.text(i, height_service + 10, f'{service_percent:.1f}%', ha='center')

plt.title('customer service vs churn')
plt.xlabel("customer service calls")
plt.ylabel('number of customers')
plt.legend(title='churn');
```

customer service vs churn






observations:


- Customers' plan are statistically significant, this means that we have an association between customer plan and churn
  - The churn rate for those with international plan is higher by about 30%
  - The churn rate for those without voice mail plan is higher by about 8%; they never get helped
- For numerical values:
  - total day minutes, total day charge, total eve minutes, total eve charge, total night minutes, total night charge, total intl minutes, total intl calls, total intl charge, customer service calls, number vmail messages are statistically significant; we reject H0
- Customer service calls vs churn
  - There is a 50% or more churn for customers who called 5 or more times

## Multivariate Analysis

In [23]:

```
corr = data.select_dtypes("number").corr()
plt.figure(figsize=(10,8))
sns.heatmap(corr, cmap='coolwarm', annot=True, fmt='.2f')
plt.title("Heatmap for numerical columns")
plt.tight_layout();
```





Observation:

- total minutes and total charges are highly correlated, this could be because minutes was used to calculate charges
  - since we saw that charges were statistically significant to churn, we will use charges as opposed to minutes

## Feature Engineering

In [24]:

```
#aggregate columns
data['total_calls'] = data['total_day_calls']+ data['total_eve_calls']+ data['total_charge'] = data['total_day_charge']+ data['total_eve_charge']
data.columns
```

Out[24]: Index(['state', 'account\_length', 'area\_code', 'international\_plan', 'voice\_mail\_plan', 'number\_vmail\_messages', 'total\_day\_minutes', 'total\_day\_calls', 'total\_day\_charge', 'total\_eve\_minutes', 'total\_eve\_calls', 'total\_eve\_charge', 'total\_night\_minutes', 'total\_night\_calls', 'total\_night\_charge', 'total\_intl\_minutes', 'total\_intl\_calls', 'total\_intl\_charge', 'customer\_service\_calls', 'churn', 'total\_calls', 'total\_charge'], dtype='object')

In [25]:

```
#drop redundant features
data = data.drop(['total_day_minutes','total_eve_minutes','total_night_
```

## Modeling

### data preprocessing

In [26]:

```
data[:3]
```

Out[26]:	account_length	area_code	international_plan	voice_mail_plan	number_vmail
0	128	415	no	yes	
1	107	415	no	yes	
2	137	415	no	no	

In [27]: `data.info(verbose=False)`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3333 entries, 0 to 3332
Columns: 17 entries, account_length to total_charge
dtypes: float64(5), int64(9), object(3)
memory usage: 442.8+ KB
```

In [28]: `#label encode the target variable`

```
le = LabelEncoder()
data['churn'] = le.fit_transform(data['churn'])
```

In [29]: `#one hot encode intl plan and voice mail plan features`

```
cat = data[['international_plan','voice_mail_plan']]
cat_ohe = pd.get_dummies(cat,drop_first=True,dtype='int')
```

`#merge the original data with the new cat columns`

```
merged_df = pd.concat([data,cat_ohe],axis=1)
```

`#drop the initial cat columns`

```
merged_df.drop(columns=cat.columns, inplace=True, axis=1)
```

```
merged_df[:3]
```

Out[29]: `account_length area_code number_vmail_messages total_day_calls total_da`

0	128	415	25	110
1	107	415	26	123
2	137	415	0	114

In [30]: `# separate features and target`

```
X = merged_df.drop("churn",axis=1)
y = merged_df.churn
```

`#split to train and test set`

```
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.2,random_state=42)
```

Out[30]: `((2666, 16), (667, 16), (2666,), (667,))`

In [31]: `#check for correlation`

```
merged_df.corr()['churn'].sort_values(ascending=False)
```

Out[31]: `churn 1.000000
international_plan_yes 0.259852
total_charge 0.231549
customer_service_calls 0.208750`

```

total_day_ccharge      0.205151
total_eve_charge       0.092786
total_intl_charge      0.068259
total_night_charge     0.035496
total_day_calls        0.018459
account_length          0.016541
total_calls             0.015807
total_eve_calls         0.009233
area_code                0.006174
total_night_calls       0.006141
total_intl_calls        -0.052844
number_vmail_messages   -0.089728
voice_mail_plan_yes     -0.102148
Name: churn, dtype: float64

```

In [32]:

```

#balance the training set
smote = SMOTE(random_state=42)
X_train_sm, y_train_sm = smote.fit_resample(X_train,y_train)
#check on whether SMOTEN worked
print(f" Original values \n {y_train.value_counts()}\n")
print(f"Smoted values \n {y_train_sm.value_counts()}")

```

```

Original values
churn
0    2280
1    386
Name: count, dtype: int64

```

```

Smoted values
churn
0    2280
1    2280
Name: count, dtype: int64

```

In [33]:

```

#scaling features
ss = StandardScaler()
X_train_s = ss.fit_transform(X_train_sm)
X_test_s = ss.transform(X_test)

```

In [34]:

```

#fit models
models = {'logistic regression':LogisticRegression(random_state=42,max_iter=100),
          'Decision tree': DecisionTreeClassifier(random_state=42),
          'Random Forest':RandomForestClassifier(random_state=42,n_estimators=10),
          'Xgboost':XGBClassifier(random_state=42,use_label_encoder=False),
          }
results = []
conf_matrix = {}
for name, model in models.items():
    model.fit(X_train_s,y_train_sm)
    y_pred = model.predict(X_test_s)
    y_prob = model.predict_proba(X_test_s)[:,1]

    results.append({'Model':name,
                    'train score':model.score(X_train_s,y_train_sm),
                    'Accuracy':accuracy_score(y_test,y_pred)*100,
                    'Recall':recall_score(y_test,y_pred)*100,
                    'Precision':precision_score(y_test,y_pred)*100,
                    'F1 score':f1_score(y_test,y_pred)*100,
                    'AUC':roc_auc_score(y_test,y_pred)*100
                    })
    conf_matrix[name] = confusion_matrix(y_test,y_pred)

```


```

results_df = pd.DataFrame(results).round(2)
results_df = results_df[["Model", "train score", "Accuracy", "Recall", "Precision", "F1 score"]]
print(results_df.to_string(index=False))

#confusion matrix
fig, axes = plt.subplots(2,2, figsize=(12,10))
axes = axes.flatten()
for i, (name, cm) in enumerate(conf_matrix.items()):
    sns.heatmap(cm, annot=True, cbar=False, cmap='Blues', fmt='d', ax=axes[i])
    axes[i].set_title(f'{name} confusion matrix')
    axes[i].set_xlabel('Predicted')
    axes[i].set_ylabel('Actual')
plt.tight_layout()
plt.show()

```

	Model	train score	Accuracy	Recall	Precision	F1 score
AUC						
71.32	logistic regression	0.79	75.11	65.98	32.49	43.54
84.31	Decision tree	1.00	85.61	82.47	50.31	62.50
83.65	Random Forest	1.00	93.25	70.10	80.95	75.14
84.87	Xgboost	1.00	88.76	79.38	58.33	67.25




Observation:

- Logistic regression misses a fair number of churners; very expensive
- Decision tree has a high recall but
- Random forest would be better if our priority was precision
- Xgboost is the best model with a balance of precision and recall (F1)

```
In [35]: #plot the roc for all models
plt.figure(figsize=(8,6))
for name,model in models.items():
    y_prob = model.predict_proba(X_test_s)[:,1]
    fpr,tpr,_ = roc_curve(y_test,y_prob)
    roc_auc = auc(fpr,tpr)
    plt.plot(fpr,tpr,label=f'{name}(AUC= {roc_auc:.2f})')

plt.title("ROC Curve Comparison")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend();
```



observations:

- Decision tree has a recall closest to our metrics of success
- So far, Xgboost is the best model to use.
- Based on our measure of success, xgboost has a recall of less than 85
- We focus on tuning it to improve the recall but first we will use cross validation

```
In [36]: #using cross validation
sm = SMOTE(random_state=42)
X_res, y_res = sm.fit_resample(X, y)

scoring = {
    'accuracy': 'accuracy',
    'recall': 'recall',
    'precision': 'precision',
    'f1': 'f1',
}

# ---- Cross-validation ----
cv = StratifiedKFold(n_splits=10, shuffle=True, random_state=42)
```

```

cv_results = {}
for name, model in models.items():
    cv_result = cross_validate(
        model,
        X_res, y_res,
        cv=cv,
        scoring=scoring,
        return_train_score=True
    )
    cv_results[name] = {
        metric: cv_result['test_' + metric].mean() * 100
        for metric in scoring.keys()
    }
    cv_results[name]['train_accuracy'] = cv_result['train_accuracy'].mean()

# ---- Results DataFrame ----
cv_results_df = pd.DataFrame(cv_results).T.round(2)
print(cv_results_df)

```

	accuracy	recall	precision	f1	train_accuracy
logistic regression	89.60	82.49	96.16	88.80	89.63
Decision tree	86.75	88.84	85.31	87.02	100.00
Random Forest	94.84	92.21	97.35	94.70	100.00
Xgboost	93.96	95.58	92.62	94.07	99.96

observation:

- we will use cross validation to split our dataset
- xgboost seems to be the best model to use
- we will hyperparameter tune to find the right parameters to use

In [37]:

```

#hyperparameter tuning
sm = SMOTE(random_state=42)
X_res, y_res = sm.fit_resample(X, y)

base_models = {
    "Logistic Regression": LogisticRegression(class_weight="balanced",
                                                random_state=42),
    "Decision Tree": DecisionTreeClassifier(random_state=42, class_weight="balanced"),
    "Random Forest": RandomForestClassifier(random_state=42, n_jobs=-1),
    "XGBoost": XGBClassifier(eval_metric="logloss", random_state=42)
}

param_grids = {
    "Logistic Regression": {
        "C": [0.01, 0.1, 1, 10],
        "penalty": ["l1", "l2"],
        "max_iter": [2000, 5000, 10000]
    },
    "Decision Tree": {
        "max_depth": [3, 5, 10],
        "min_samples_split": [5, 10, 20],
        "min_samples_leaf": [2, 5, 10]
    },
    "Random Forest": {
        "n_estimators": [100, 200],
        "max_depth": [5, 10, 15],
        "min_samples_split": [5, 10],
        "min_samples_leaf": [2, 5]
    },
    "XGBoost": {
        "n_estimators": [100, 200],
        "max_depth": [5, 10, 15]
    }
}

```

```

    "max_depth": [5, 7, 10],
    "learning_rate": [0.05, 0.1, 0.2],
    "subsample": [0.8, 1],
    "colsample_bytree": [0.8, 1],
    "reg_lambda": [1, 5],
    "reg_alpha": [0, 1]
}
}

scoring = {
    "accuracy": "accuracy",
    "recall": "recall",
    "precision": "precision",
    "f1": "f1",
    "roc_auc": "roc_auc",
    "pr_auc": "average_precision",
}

cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

best_models = {}
cv_results = {}

for name, model in base_models.items():

    grid = GridSearchCV(
        model,
        param_grids[name],
        cv=cv,
        scoring="recall",
        n_jobs=-1,
        verbose=0
    )
    grid.fit(X_res, y_res)
    best_model = grid.best_estimator_

    if name == "Decision Tree":
        best_model = CalibratedClassifierCV(best_model, cv=5)
        best_model.fit(X_res, y_res)

    best_models[name] = best_model

    cv_result = cross_validate(
        best_model,
        X_res, y_res,
        cv=cv,
        scoring=scoring,
        return_train_score=True
    )
    cv_results[name] = {
        metric: cv_result["test_" + metric].mean() * 100
        for metric in scoring.keys()
    }
    cv_results[name]["train_accuracy"] = cv_result["train_accuracy"].mean()
    cv_results[name]["best_params"] = grid.best_params_

    cv_results_df = pd.DataFrame(cv_results).T
    cv_results_df = cv_results_df.applymap(lambda x: round(x, 2) if isinstance(x, float) else x)

print(cv_results_df)

```

```

/opt/anaconda3/envs/learn-env/lib/python3.9/site-packages/sklearn/linear
_model/_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (s
tatus=1):
STOP: TOTAL NO. OF ITERATIONS REACHED | TMIT

```

```
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result()
```

```
/opt/anaconda3/envs/learn-env/lib/python3.9/site-packages/sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status=1):
```

STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result()
```

```
/opt/anaconda3/envs/learn-env/lib/python3.9/site-packages/sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status=1):
```

STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result()
```

```
/opt/anaconda3/envs/learn-env/lib/python3.9/site-packages/sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status=1):
```

STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result()
```

```
/opt/anaconda3/envs/learn-env/lib/python3.9/site-packages/sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status=1):
```

STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result()
```

```
/opt/anaconda3/envs/learn-env/lib/python3.9/site-packages/sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status=1):
```

STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result()
/opt/anaconda3/envs/learn-env/lib/python3.9/site-packages/sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status=1):
```

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result()
/opt/anaconda3/envs/learn-env/lib/python3.9/site-packages/sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status=1):
```

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result()
/opt/anaconda3/envs/learn-env/lib/python3.9/site-packages/sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status=1):
```

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result()
/opt/anaconda3/envs/learn-env/lib/python3.9/site-packages/sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status=1):
```

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result()
/opt/anaconda3/envs/learn-env/lib/python3.9/site-packages/sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status=1):
```

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result()
```

```
/opt/anaconda3/envs/learn-env/lib/python3.9/site-packages/sklearn/linear
```

```
/opt/anaconda3/envs/learn-env/lib/python3.9/site-packages/sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result()
```

```
/opt/anaconda3/envs/learn-env/lib/python3.9/site-packages/sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status=1):
```

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result()
```

```
/opt/anaconda3/envs/learn-env/lib/python3.9/site-packages/sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status=1):
```

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result()
```

```
/opt/anaconda3/envs/learn-env/lib/python3.9/site-packages/sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status=1):
```

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result()
```

```
/opt/anaconda3/envs/learn-env/lib/python3.9/site-packages/sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status=1):
```

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result()
```

```
/opt/anaconda3/envs/learn-env/lib/python3.9/site-packages/sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status=1):
```

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result()
/opt/anaconda3/envs/learn-env/lib/python3.9/site-packages/sklearn/linear_
model/_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (s
tatus=1):
```

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result()
/opt/anaconda3/envs/learn-env/lib/python3.9/site-packages/sklearn/linear_
model/_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (s
tatus=1):
```

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result()
/opt/anaconda3/envs/learn-env/lib/python3.9/site-packages/sklearn/linear_
model/_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (s
tatus=1):
```

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result()
/opt/anaconda3/envs/learn-env/lib/python3.9/site-packages/sklearn/linear_
model/_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (s
tatus=1):
```

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result()
/opt/anaconda3/envs/learn-env/lib/python3.9/site-packages/sklearn/linear_
model/_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (s
tatus=1):
```

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```

egression
    n_iter_i = _check_optimize_result(
        accuracy  recall  precision     f1  roc_auc  pr_auc
    \
    Logistic Regression   89.56   83.33    95.20  88.87   95.43  96.54
    Decision Tree         87.91   80.67    94.34  86.97   94.30  95.55
    Random Forest         92.74   87.86    97.36  92.36   98.10  98.40
    XGBoost               95.05   93.26    96.73  94.96   98.58  98.84

        train_accuracy \
Logistic Regression      89.67
Decision Tree            91.58
Random Forest            98.88
XGBoost                  100.00

        best_params
Logistic Regression      {'C': 0.1, 'max_iter': 5000, 'penalty': 'l2'}
Decision Tree             {'max_depth': 10, 'min_samples_leaf': 5, 'min_...
Random Forest             {'max_depth': 15, 'min_samples_leaf': 2, 'min_...
XGBoost                   {'colsample_bytree': 1, 'learning_rate': 0.2, ...}

observation: Model train score Accuracy Recall Precision F1 score AUC logistic
regression 0.79 75.11 65.98 32.49 43.54 71.32 Decision tree 1.00 85.61 82.47 50.31
62.50 84.31 Random Forest 1.00 93.25 70.10 80.95 75.14 83.65 Xgboost 1.00
88.76 79.38 58.33 67.25 84.87



- Imbalance handled with SMOTE; recall boosted from 65% (Logistic baseline) to >90% (XGBoost).
- Overfitting addressed; Decision Tree pruned with parameters
- from feature importance: international plan, customer service calls, and high day charges.
- Business Value: With XGBoost, we can now identify >93% of churners, directly supporting customer retention strategies.

```

In [38]:

```

#Feature Importance using XGBoost
best_xgb = best_models["XGBoost"]

xgb_importances = best_xgb.feature_importances_
feat_imp = pd.DataFrame({
    "Feature": X.columns,
    "Importance": xgb_importances
}).sort_values(by="Importance", ascending=False)

plt.figure(figsize=(10,6))
sns.barplot(x="Importance", y="Feature", data=feat_imp.head(15), palette="Blues")
plt.title("Top 15 Feature Importances (XGBoost - Gain Based)")
plt.tight_layout()
plt.show()

```

